

PRODUCT OVERVIEW

Real interviews. Real signal. Ten times faster.

Roundz is the AI interview engine for the AI-native era.

Roundz is a platform built around one thing: an **AI Interviewer** that takes your complete technical loop end to end. It conducts every round by voice the way a senior engineer would, asking, listening, probing the reasoning, and watching the candidate code and design as they build. Each round is scored against your bar and returned as one explainable report. Hundreds of applicants become a ranked shortlist before a senior engineer spends an hour on first-round screens, and your team still makes every final call.

IN ONE LINE

The AI that runs the whole interview, so your team only meets the best.

Hiring engineers is slow and noisy: hundreds of applicants per role, and a static coding test that tells you an answer arrived without telling you how. The role has moved too. Engineers now build alongside AI assistants every day, so what separates them is reasoning, decomposition, and judgment rather than recall. Roundz measures candidates in that environment instead of the traditional one. Its **AI Interviewer** conducts a real multi-round loop by voice, watches the code and system design as they are built, and probes the thinking behind every answer. Where a team also wants to see AI fluency, an optional round gives the candidate an **AI Copilot** and scores how well they direct it. Everything is measured against your bar and returned as one explainable report. Roundz filters; your team decides.

CONTENTS

01	Why the technical screen stopped working	03	08	One dashboard: source, screen, shortlist	10
02	What Roundz is, and who it is for	04	09	Capability and what it means for you	11
03	Two agents, two different jobs	05	10	Trust, fairness, and the limits we hold	12
04	The agentic-coding round	06	11	Commercial model and the one-req pilot	13
05	The whole loop, not a single test	07	12	Glossary and reusable one-liners	14
06	How an interview runs, step by step	08	13	Your first req is on us	15
07	The report your team receives	09			

WHO SHOULD READ WHAT

EXEC OR FOUNDER

Sections 01 to 04: what changed in the role, what the product is, the agents that run it, and why it exists now. Section 11 covers how a first engagement starts.

HIRING MANAGER

Sections 05 to 08: the loop, how an interview actually runs, the report you receive, and the pipeline you run it from.

TECHNICAL BUYER

Sections 09 and 10: the full capability list, and the limits we deliberately hold ourselves to.

01 / THE PROBLEM

Why the technical screen stopped working

Two things are true at once: a static test cannot show you how a candidate thinks, and the job it was written to test for has changed.

SHIFT ONE

A test cannot see how someone thinks

A coding test tells you an answer arrived. It cannot show you how the candidate framed the problem, what they ruled out and why, how they held up when pushed on a trade-off, or whether they could defend the design out loud. Those only surface in conversation. It is why candidates routinely clear an automated screen and then fail the first live round: the two are measuring completely different things, and only one of them resembles the job.

SHIFT TWO

The job itself changed

Engineers now spend their day directing AI tools, not typing every line. The skill that decides output quality is judgment: prompting, decomposition, and catching the model when it is wrong. Testing someone on the thing they no longer spend their day doing tells you very little about how they will perform in the role you are hiring for.

THE NET EFFECT

Teams are running static filters for a dynamic role. Candidates clear the automated screen, then fail the live round because they cannot explain their logic or defend a trade-off. Each of those is an hour taken from a senior engineer who could have been shipping, spent discovering something the screen should have caught, and it repeats for every candidate the filter lets through.

WHAT THE STATUS QUO COSTS

- A poor hire can cost up to ~30% of first-year salary.

US Dept. of Labor / SHRM

- Interviewing drains your scarcest resource: senior engineer time.

Stripe Developer Coefficient

- Turnover and ramp costs from mis-hires compound quarter over quarter.

Officevibe / LinkedIn

- AI-assisted cheating on remote screens is documented and rising.

Glider AI

Figures above are third-party industry research shown for context. They are not Roundz performance claims.

02 / THE PRODUCT

What Roundz is, and who it is for

An AI that conducts the complete interview loop, so the people who reach your calendar have already cleared a real, consistent bar.

Roundz's AI Interviewer runs each round the way a strong human interviewer would: it talks, listens, asks spoken follow-ups, and probes the reasoning behind every answer rather than the final output. It watches the candidate code and design in real time, scores every round against your rubric, and hands your team an explainable report. The experience is built to feel like a real, face-to-face interview.

Because the same loop, the same rubric, and the same bar run for every candidate on a role, the comparison between two people is finally apples to apples. That consistency is what makes the shortlist worth trusting, and it is the part manual screening has never been able to guarantee.

THE 80/20 MODEL



AUGMENTATION, NOT AUTOMATION

Roundz does the heavy, repetitive interviewing at a consistent bar and surfaces the roughly top 20% worth a human's time. It runs the full loop and scores it, then stops precisely where judgment begins: its job is to work out which candidates are genuinely worth your final rounds, and to hand your team the evidence behind that shortlist.

A TWO-SIDED PRODUCT

Employers and hiring teams

Design a full interview loop once, let Roundz run every round at scale, and review explainable, auto-scored reports. Senior engineers spend their time only on pre-vetted candidates, and every report is auditable if a decision is ever questioned. This is where the commercial focus sits.

Candidates

Practice realistic, full-loop technical interviews with the AI Interviewer (Technical, design, behavioral, agentic coding) and get honest, detailed feedback on how they did, scored to the level they are targeting. It also feeds the employer side: a growing pool of engineers who have already completed a full loop and carry a scored, verifiable read, so some reqs can be filled from candidates Roundz has evaluated already.

Two agents, two different jobs

Roundz runs on one agent: the AI Interviewer, which conducts every round. A second agent, the AI Copilot, is optional, switch it on to see how a candidate works with AI and leave it off to run a traditional interview.

● ALWAYS ON, WORKS FOR YOU

The AI Interviewer

Takes the interview the way one of your senior engineers would: it talks and listens in real time, asks the spoken follow-up, pushes on a trade-off, and watches the code and design as they are written.

● OPTIONAL, YOUR CALL

The AI Copilot

A coding assistant you can choose to give the candidate, in the one round designed around it. Enable it when AI fluency is part of the bar; leave it off and that round never appears. When on, it genuinely helps them, exactly like the copilot on their desk at work.

	THE INTERVIEWER	THE COPILOT
WHOSE SIDE	Yours. It carries your bar into every conversation.	The candidate's, when you choose to enable it.
ITS JOB	Conduct the round, probe the reasoning, score it.	Help the candidate decompose and build.
DOES IT HELP?	Only as a real interviewer would: a scoped hint when someone stalls, and it is recorded.	Yes, fully. Helping is the entire point of giving it to them.
WHERE IT APPEARS	Every round of every loop, start to finish.	Off unless you turn it on. One optional round, added when AI fluency matters.
WHAT IT PRODUCES	The evaluation, the verdict, and the report.	A measurable read on AI fluency: problem decomposition, prompt quality and sequence, attempts taken, what they accepted or corrected, and token cost.

THE LINE WE NEVER CROSS

The Interviewer never writes the candidate's code, and the Copilot never scores anyone. That separation is what makes the optional round a real measurement instead of a favour: the candidate gets genuine help, and Roundz still learns exactly how good they are at asking for it.

The agentic-coding round

In this one round the candidate is given the AI Copilot and asked to solve a problem with it. The Interviewer keeps running the round; what it now scores is how well the candidate directs the Copilot. No other interview platform measures this.

Every coding test ever built answers one question: can this person write correct code by hand? AI tools now write correct code for almost anyone, so that question no longer separates strong engineers from weak ones. Roundz asks the question that does. The Copilot the candidate gets is a real one, with chat and inline completions, and it genuinely helps them, which is the point. Roundz records the full trace of how they used it: every prompt, every accepted, rejected or edited suggestion, and how efficiently they spent a fixed token budget. The Interviewer evaluates that trace alongside everything else it observed.

01 PROMPTING

Do they ask the AI for the right thing, clearly, with the context it needs? Vague prompts produce plausible code that solves the wrong problem, and it shows.

02 DECOMPOSITION

Do they break a hard problem into the right sub-problems before delegating? Strong engineers scope the work first; weaker ones hand over the whole thing and hope.

03 JUDGMENT

Do they notice when the AI is wrong and correct it, or accept bad output? This is the single clearest separator between candidates in the round.

04 EFFICIENCY

Do they reach a good answer without burning effort, context, or budget? Getting there in three good turns is a different skill from getting there in thirty.

INSTRUMENTED

Prompt text, sub-prompt structure, accept / reject / edit events, code-context deltas, latency.

BOUNDED

A hard per-session token budget, so token efficiency becomes part of the signal, not an afterthought.

SEPARATELY SCORED

A dedicated evaluation path produces a competency assessment for AI-assisted work.

WHY NOW

Two things arrived together: voice models became good enough to hold a real technical conversation, and building alongside an AI assistant became the normal way engineers work. The interview had to change, and only now do the tools exist to change it.

The whole loop, not a single test

Most tools automate one screen. Roundz conducts the entire multi-round loop a candidate would face in real life, each round live, voice-led, and auto-scored.

Build it once

Assemble a loop from a library of ready-made rounds, set the order, duration, and level, then reuse it across every role that shares that bar.

Calibrated to level

Every round is tuned to entry, mid, or senior expectations, with question difficulty and rubric weights that move with it.

Rolls up cleanly

Each completed loop consolidates into strengths, growth areas, and one clear hire signal, per round and overall.

Interview Loops

ROUNDZ.AI

Sequences of rounds, reusable across roles

+ Create New Loop

LOOP NAME	ROUNDS	DURATION	COMPLETION	USAGE
Full Stack Engineering Senior, 3-round loop	3 rounds	175 min	7 days	24 times
Data Engineering Mid-level screening loop	2 rounds	115 min	7 days	18 times
Software Development Engineer Entry-level, DSA focused	2 rounds	90 min	5 days	31 times
Cloud Engineer Design-weighted loop	3 rounds	150 min	5 days	9 times
Marketing Executive Non-technical behavioral loop	1 rounds	45 min	3 days	4 times

The loop builder: reusable loops with their round count, total duration, and completion window. Illustrative data.

06 / THE MECHANICS

How an interview runs, step by step

From "a candidate applies" to "your team decides."

- 1 The role and loop are set up.** The employer posts a job and assembles the interview loop. Roundz extracts the skills the role needs and ranks inbound applicants by the meaning of their experience, not keyword overlap.
- 2 The live voice interview runs, round by round.** The AI Interviewer sets up the problem, listens while the candidate works through it, and asks what a human would ask next: why that approach, what breaks when this scales, what you would drop if you had half the time. Candidates can cut in and think out loud, and it waits instead of talking over them.
- 3 Roundz watches the candidate work.** The code editor and the design whiteboard stream to the agent as the candidate builds, so it reasons about the evolving solution rather than a final artifact. Diagrams are read structurally, not as pictures.
- 4 Phases advance on their own timers.** Each round is time-boxed and the interviewer moves on cleanly. It is deliberately tuned not to abandon a candidate who has gone quiet mid-problem, and it survives a dropped connection.
- 5 Every round is scored against a rubric.** Weighted criteria keyed to the role's level and the question type, applied identically to every candidate for that role.
- 6 The human decides.** Roundz has run the loop and done the filtering. Your team reads the report and makes the call.


Technical Interview - Software Engineer

54:32

PROBLEM STATEMENT A 0-indexed array derived is formed by the bitwise XOR of adjacent values in a binary array original. Determine whether a valid original could have formed it. Input: derived = [1,1,0] Output: true	Language JAVA <pre> 1 class Solution { 2 public boolean doesValidArrayExist(int[] derived) { 3 int xor = 0; 4 5 for (int n : derived) { 6 xor ^= n; 7 } 8 9 return xor == 0; 10 } 11 }</pre>	TRANSCRIPT INTERVIEWER Walk me through the reasoning behind this solution, and why XOR is the right operation here. CANDIDATE Each element is the XOR of two adjacent originals, so folding the whole array cancels every value twice. INTERVIEWER Good. What breaks if the array were not binary?
INTERVIEWER SPEAKING 		

The candidate's view during a coding round: the problem, a live editor the Interviewer reads as it is typed, the round timer, and the spoken exchange captured as a transcript. Illustrative data.

07 / THE OUTPUT

The report your team receives

Not a mysterious number. A structured, auditable read on one candidate that a hiring manager can defend line by line.

Technical and behavioral narrative

What the candidate demonstrated in each round, written out, with the reasoning behind the read.

Strengths and growth areas

Individually itemised, so a debrief starts from evidence rather than impressions.

Per-criterion scores

Each rubric criterion scored and weighted, with the level (entry, mid, senior) it was judged against.

Match and integrity context

The semantic resume-to-role match score, matched and missing skills, and any proctoring signals raised.

THE VERDICT SCALE

STRONG_HIRE

HIRE

LEANING_HIRE

NO_HIRE

STRONG_NO_HIRE


Full Stack Engineering · Candidate Report

ROUNDZ.AI

EW

E. Walker
Full Stack Engineer · Mid-level · 4 yrs exp

VERDICT
HIRE

LOOP PROGRESS

3 / 3 rounds

✓

DS & Algorithm Coding

8.4

✓

Low Level Design

7.9

✓

Full Stack Engineering

8.1

SIGNAL

80%

role match

Strengths
Clear complexity reasoning; defended the caching trade-off unprompted.


Growth
Shallow on failure modes at scale.

A completed loop: every round scored, the role match, itemised strengths and growth areas, and the verdict your team acts on. Illustrative data.


08 / THE PIPELINE

One dashboard: source, screen, shortlist

Roundz is not a bolt-on test. It is the hiring pipeline for the technical funnel, from an open req to a ranked shortlist.



Every stage is visible on the board, so you always know where a requisition actually stands and which candidates are waiting on a decision.



The employer dashboard: every job rack with the stage its pipeline has reached. Illustrative data.

09 / CAPABILITIES

Capability, and what it means for you

Every capability below is shipped and in the product today.

CAPABILITY	WHAT IT MEANS FOR YOU
The AI Interviewer, by voice	A real conversation with human-like turn-taking, not a static test. It asks the follow-up and pushes on the trade-off, at scale, around the clock, with no scheduling.
Live code and whiteboard awareness	The agent sees the editor and the design canvas as the candidate builds, and judges approach, complexity, and correctness of reasoning rather than a final artifact.
AI Copilot for candidates (optional)	Enable it per loop when AI fluency matters for the role. The candidate gets a real copilot that helps them, and Roundz scores how they directed it: prompt quality, problem decomposition, judgment on bad output, and token efficiency. Leave it off to run a fully traditional interview.
Multi-round configurable loops	Design the loop once from ready-made round types, set order and duration, and reuse it across roles so the bar stays identical.
Design your own interview	Build the loop's building blocks yourself: author custom rounds, add your own questions to the knowledge bank with their reference answers, and define the evaluation rubric each round is scored against.
Explainable, level-aware reports	Rubric-scored to entry, mid, or senior expectations with a verdict and the evidence behind it. Defensible to your team and to a candidate who asks why.
Anti-cheat proctoring	Gaze and face signals plus keyboard, visibility, and clipboard telemetry flag integrity risk, so the shortlist you act on is trustworthy.
Semantic resume to role matching	Inbound applicants ranked by real fit rather than keyword overlap, with matched and missing skills surfaced per candidate.
Curated question and knowledge bank	Questions carry reference answers, progressive hints, and starter signatures across 20+ languages, mapped to the concepts each round is meant to test.
Interview reliability engineering	Reconnection grace periods and a tested decision policy so the interviewer does not abandon a candidate who goes quiet while thinking.

Trust, fairness, and the limits we hold

Being precise about what an AI does and does not do is how you earn a technical buyer's trust.

- **Explainable by design.** Every score carries its reasoning and maps to a defined rubric and the role's level. That matters for good decisions, and for a rising regulatory bar around AI in hiring that increasingly expects transparency and human oversight.
- **A consistent bar.** Every candidate for a role runs the same loop against the same criteria, which removes the interviewer-to-interviewer variance that makes manual screening noisy.
- **Integrity signals, not accusations.** Gaze, face, keyboard, visibility, and clipboard telemetry are surfaced as flags for a human to interpret. Roundz raises the signal; it does not disqualify anyone.
- **Human-in-the-loop, always.** The final call is a person's. That is a design constraint, not a disclaimer.

WHAT ROUNDZ DOES NOT DO

It does not compile or execute the candidate's code.

This is deliberate. The round is built to work like a face-to-face interview, where you watch someone reason through a problem and talk it through with them, not like an automated judge waiting on test cases. Roundz evaluates approach, complexity, and the correctness of the reasoning from the code as it is written and from reference answers, exactly as your own interviewer would sitting beside them.

It does not make the hiring decision.

Roundz runs the loop, scores it, and tells you which candidates are worth your final rounds and why. Acting on that shortlist stays with your team.

It is not a resume keyword filter.

Matching reads the meaning of a candidate's skills, and the real signal comes from the live interview loop, not the resume.

It does not replace your final rounds.

Roundz is built for the first-round technical loop. Culture, team fit, and the closing conversation stay where they belong, with your people, on the candidates Roundz has already qualified.

11 / COMMERCIAL

How you buy it, and how you start

Billing meters the one thing you actually consume: a candidate interview.

PILOT PACK

One requisition

A prepaid block of credits for a single open role, plus rubric calibration, weekly shortlists, and an end-of-pilot readout.

GROWTH PACKS

Buy credits in blocks

For spiky or low-commitment hiring. Larger blocks carry a lower per-credit price. All rounds included.

ANNUAL

Committed volume

Lowest effective per-credit price, seats, priority support, and at higher tiers SSO, ATS integration, and custom rubric calibration.

ONE CREDIT, ONE CANDIDATE INTERVIEW

Consumption is tracked in a full audit ledger, so you can see exactly what every credit was spent on.

THE ONE-REQ PILOT, AND HOW WE JUDGE IT

Pick one role you are actively hiring for. We run your real loop on the next cohort of candidates, calibrated to your bar with one of your engineers at kickoff, side by side with your current screen. Success criteria are agreed in writing before we start.

01 A measurable cut in the senior-engineer hours spent on first-round screens for that role.

02 Shortlists delivered within the week a candidate interviews, fast enough to act on live pipeline.

03 Hiring-manager agreement with the Roundz verdict on the shortlist, measured against their own calls.

04 No material increase in candidate drop-off compared with the current screen.

WHAT YOU BRING

One active req and its target level, access to the candidate flow for it, and roughly an hour of one engineer's time at kickoff to calibrate the rubric to your bar.

WHAT YOU GET BACK

An explainable report on every candidate screened, weekly shortlists, a side-by-side read against your current process, and a joint readout against the criteria above.

Glossary and reusable one-liners

TERM	PLAIN MEANING
The AI Interviewer	Roundz's voice agent that conducts the interview on your behalf, in every round of every loop. It probes and scores; it never writes the candidate's code.
The AI Copilot	An optional assistant you can give the candidate in the round built for it. It genuinely helps them; how they drive it is what gets scored. Leave it off for a traditional interview.
Interview loop	The full sequence of rounds a candidate goes through. Roundz runs the whole loop, not one screen.
Round	One interview inside the loop, for example a coding round or a system-design round.
Agentic coding	Solving a problem by directing an AI assistant rather than writing every line by hand. Roundz has one optional round for it, using the AI Copilot.
80/20, human-in-the-loop	Roundz does the roughly 80% that is screening work; your team keeps the 20% that is judgment and the decision.
Rubric, level-aware	A defined scoring guide of weighted criteria. Scoring adjusts to whether the role is entry, mid, or senior.
Semantic matching	Matching a resume to a job by the meaning of the skills involved, not keyword overlap.
Proctoring	Integrity monitoring during an interview, surfaced as signals for a human to interpret.
Credit	The billing unit: one credit covers one candidate interview, tracked in an audit ledger.
Verdict	The overall hire recommendation, from strong hire to strong no-hire, always with its reasoning.

ONE-LINE SUMMARIES YOU CAN REUSE

FOR AN EMPLOYER

"Roundz's AI Interviewer runs your entire technical loop by voice, coding, design, behavioral, and an agentic round where the candidate is given an AI Copilot and scored on how well they direct it. Your team gets explainable, auto-scored reports, meets only pre-vetted candidates, and still makes every final call."

FOR A CANDIDATE

"The first interview that measures how well you direct AI. Practice a full loop with the AI Interviewer, solve a round alongside an AI Copilot, and get an honest, level-aware read."

FOR ANYONE

"Roundz is an AI that conducts complete technical interviews, round by round, and feels like the real face-to-face thing. In one round it hands the candidate an AI copilot and measures how well they use it."

13 / NEXT STEP

Your first Job Req is on us.

Pick one role you are actively hiring for and we will run your real interview loop on it, free. No credits, no commitment. We agree what success looks like before we start, you run it side by side with your current screen, and you judge the reports against your own calls. If they earn your team's trust, we go from there.

 WEBSITE

roundz.ai
 BOOK A DEMO

roundz.ai/employer
 EMAIL

support@roundz.ai


WHAT HAPPENS NEXT

01 PICK THE ROLE

A 30 minute call to choose the requisition and confirm the level you are hiring at. Nothing to sign.

02 CALIBRATE

One of your engineers tunes the rubric to your bar, so the first report already speaks your language.

03 RUN IT LIVE

Your real candidates, your real loop, with shortlists back in the same week they interview.

IN ONE LINE

The AI that runs the whole interview, so your team only meets the best.